



ORANGUTAN ENCYCLOPEDIA | ATTEMPTER GUIDELINES

CUA v3 Task Creation Guidelines for Attempters

Version	Changes
V1.0	Initial release
V2.0 06/05	Updated Workflow Video + New Task Instructions • Updates about the VM. • New "Execute Task Initializer" step in VM • Preferred workflow: Work inside the VM • File Naming Requirements • New final step: Agent Run • General improvements
V2.01 06/11	Added clarifications about what the prompt must include: - The exact name of the output files is supposed to be. (in a natural way). - The location they must be saved in: All output files must be saved on the desktop, and the prompt must state this.

Version	Changes
	- Remove Personally Identifiable Information (PII) in the prompt content.
V2.02 06/15	- Added the current issue with expiring URLs and how to solve it.
V2.03 06/18	- Replaced “prompt must have 3 large attachments” in section 3.3 with “Initial files should be a sufficient baseline for a complex prompt.”
V3 6/25	• Prompt: Incorporated GUI interactions • Rubric: Will be autogenerated + Updated guidelines. • Computer Use (VM): Clarifications on pasting URLs & filenames into the initializer + verifier • Agent Run: goal is to achieve an average score $\leq 65\%$ ○ Approximately ~20-30 min to run
V3.01 6/26	Clarification: • Agent run step takes approximately 20-30 min to generate responses • Prefilled tasks: You may change the number and/or file format of the pre-seeded output files (initial output files) • Updated verifier JSON example Added Rubric Generator
6/29	Added a ‘Helpful Resources’ section for easy access Clarification: • While the autogenerated rubric assigns weights that add up to 100, contributors may add/delete and/or adjust the weights as long as the weight share follows the target weight share in section 5.2 • The rubric generator no longer requires input files to be uploaded. You only need to paste the prompt and GTF files. • Clarifications in the initializer & verifier JSON section • If the VM desktop does not contain the GTF files, the verifier will not return a 1. ○ If the GTF files are in the VM but the verifier doesn’t score a 1, you may proceed with the task.

Version	Changes
	If you encounter issues in the agent run step (e.g., the agent fails to run), please provide feedback at the end of the task, and you may submit the task.
6/30	Rubric Generator: - You may use the URL in the Asset Upload section directly in the rubric generator. You no longer have to convert the URL formats. - Input files are required as input Rubric Category: 'Format Gates' is reserved for
7/1	Visual Criteria in Rubric: Implicit requests in the prompt, or situations where there can be multiple correct interpretations to visual elements, do not need to anchor to the GTF (expected file).
7/2	Clarifications: If contributors would like to use the initial output files (pre-seeded output files) but they are not in the VM when you execute the initializer, please transfer them to VM yourself (e.g., via filebin, Google Drive, etc).
7/14	Added Rubric Examples doc to Helpful Resources and Rubric Section (section 5)
7/21	Updated rules in the rubric section
7/29	New step: GTF eval check. You will see a GTF eval in the next step. It is designed to help sanity check your GTF files.

Version	Changes
8/1	Agent Run removed from the attempter flow. CBs no longer run the agent run step. Added instructions to avoid reward hacking + new information to avoid common errors in QC
8/18	Rubric (GT Anchored sub-section) - Add cross file consistency examples. Cross file consistency criteria do not have to anchor to the GT Rubric (Weight Share) - Added a different set of weight share for Design & Creative domain Prompt (Section 2.3 E) - If the prompt requires work done in the Blender app, please design the task with a .blend output file instead of asking for screenshots of the 3D structure in a pdf, docx, etc. - You may want to include a .blend input file and ask the agent to make modifications to the input file and provide the resulting .blend file as an output.
8/26	Inclusion of copilot in task Added new parameters for Ddesign and Multimedia tasks

Confidentiality Notice

This document is for internal use by authorized Orangutan_Encyclopedia attempters only.
Do not share externally or post in public forums.

Helpful Resources



[Outlier Community](#)



[Task Workflow Video](#)



[Rubric Generator Video](#)



[FAQ and Common Errors](#)



[Prompt Writing Handbook](#)



[Rubric Examples doc](#)

Table of Contents

[Helpful Resources](#)

[1. Overview](#)

[2. Prompt section](#)

[Quick Reference: Prompt Editing Decision Flow](#)

[3. Skip Asset Upload & Rubric for now](#)

[4. Computer Use Step. VM.](#)

[4.1 Task Initializer](#)

[4.2 Create GTF Files](#)

[4.3 Upload GTF Files to the Asset Section](#)

[5. Rubric Section](#)

[5.1 Overview](#)

[5.2 Scoring Model](#)

[5.3 Write Criteria Against the Reference](#)

[5.4 Tolerance Rules](#)

[5.5 What Not To Do](#)

[6. Verifier](#)

[7. GTF eval check](#)

1. Overview

Orangutan Encyclopedia is a **Computer Use Agent (CUA) benchmark project** for training frontier AI models on real-world desktop and browser automation tasks.

These guidelines outline the task creation flow. Your task is to create high-quality, verifiable tasks across numerous professional domains.

The AI model will solve the prompt and create its output file, which will then be scored against a rubric based on your Ground Truth Files (GTF). The goal is to achieve model failure.

1.1 Glossary


Term	Definition
GTF/ ground truth file	The output file created by contributors
Expected file	The output file created by contributors. Use 'expected file' in the rubrics instead of 'GTF/ground truth file'.
Initial Output files	Pre-seeded output file in prefilled tasks. Previously referred to as 'gold files'.
Virtual Machine (VM)	A computer running inside another computer. It gives you a separate environment where you can safely open, edit, and save files without affecting your own system.
Graphical User Interface (GUI)	A visual way to interact with a computer or device. Instead of typing text commands, a GUI lets you control the device by clicking or tapping on pictures, icons, and buttons.

1.2 Task workflow overview

You will encounter 2 types of tasks:

1. empty tasks to start from scratch; or
2. prefilled tasks with a prompt + input files
 - a. You may edit/replace the input files
 - b. You may change the number and/or format of the initial output files pre-seeded with the task.

The only difference is that blank tasks have an extra sourcing step to find input files. The rest of the steps are the same for all tasks.



For reference on the overall task workflow and taxonomy, **please watch this video!**

[TASK WORKFLOW VIDEO](#) - Click here.

It will help you a lot in understanding the overall context and getting up to speed!

This video is also attached to the task taxonomy, so you can refer back to it whenever you have questions or need a refresher on the workflow.

1.3 Identify Task Domain

Before you start with your prompt, check the assigned domain and subdomains at the top of the task. The task domain and subdomain **must match your expertise and skill set**.

- If you encounter a task that is outside your domain or an area you're not comfortable reviewing, feel free to skip it and move on to another task.
- If the task has a prefilled prompt and/or input files about a different topic than the subdomain assigned to the task, replace the prompt and/or input files.

Welcome to the Orangutan Encyclopedia Project!

Task Domain

Parent Domain: [domain placeholder]

Sub Domain: [sub_domain placeholder]

Your main work as an Attempter has three parts:

1. Review a provided prompt ensuring that it is natural, feasible and follows all project guidelines
2. Review the provided rubrics and ensure that they are a valid rubric based on project guidelines. If the rubric has any issues, you must fix
3. Upload the gold files (answers) into the VM and verify that the gold files correctly answer the prompt based on the rubric/your own evaluation.
4. Run the verifier and ensure that the gold file gets a perfect score.

Important Note: If you see a computer environment error, refresh your browser window.
Your task progress is saved automatically - you will **not** lose your work

See this video for a live walkthrough: <https://static.remotasks.com/uploads/Onboarding video sped.mp4>

Domains	Description
Finance & Accounting	Accountants, Auditors, Financial Managers, Financial Analysts, Personal Advisors, Securities Agents
Software & Data	Software Developers, IT Managers, Data Scientists, Game Designers
Design & Creative	Graphic Designers, UI/UX Designers, 3D Modelers
Engineering	Industrial Engineers, Mechanical Engineers, Chemical Engineers, CAD Designers, Architects
Legal & Compliance	Lawyers, Compliance Officers, Paralegals
Healthcare	Registered Nurses, Nurse Practitioners, Pharmacists, Medical Admins, Health Service Managers
Sales & Marketing	Sales Managers, Sales Reps, Customer Service Reps, Retail Supervisors
STEM Research	Biologists, Chemists, Physicists, Research Scientists, Bioengineers
Multimedia & A/V	Producers, Film Editors, A/V Technicians, Journalists
Operations & Management	Admin Managers, Shipping Clerks, Project Mgmt Specialists, Operations Managers, Management Consultants
Real Estate	RE Brokers, RE Sales Agents, Property Managers, Buyers/Purchasing Agents
Government & Public Sector	Concierges, Recreation Workers, Social Workers, Detectives

2. Prompt section

For Full Domain and Application Specific Prompt Guidance, See the [Prompt Writing Handbook](#)

Purpose of this section

A strong prompt gives the agent one clear interpretation of the task. It explains what to produce, which files to use, where to save the output, and what success looks like without sounding like a grading rubric. The final prompt should feel like a real workplace request while still being specific enough to support reliable GTF files, rubrics, and verification.

This section is organized into three main parts:

1. Review the task context and/or source input files.
2. Rewrite/Create the prompt.
3. Validate and finalize the prompt.

Caution for the Healthcare and Biology domains

Avoid prompts that deal with topics like **influenza, COVID, Hantavirus**, and other highly contagious pathogens of that nature - the agent will refuse to run this due to internal guardrails dealing with biorisk.

It doesn't matter if the prompt's context is entirely innocent!

2.1 Review the Task Context (for prefilled tasks)

Before editing the prompt, understand what the task is supposed to test and what materials the agent will have available. Do not assume the seeded prompt or initial output files are already correct.

Item to review	What to check	What to do if there is an issue
Assigned domain and subdomain	Confirm that the prompt matches the domain shown in the goal section.	If the prompt/input files clearly belong to a different domain, replace them.
Primary prompt	Focus on the prompt shown on the left side of the task page.	Edit this prompt when it is unclear, unnatural, too easy, or misaligned with the files. Also add GUI requests.

Item to review	What to check	What to do if there is an issue
Starting files	Confirm that the input files are usable, supported, stable, and sufficient for the requested task.	If the source files contain PII, are corrupted, unsupported, or insufficient, edit and/or replace them.
Initial output files	Compare the Initial output files against the prompt and input files.	Correct/Replace the Initial output files if they do not fully satisfy the final prompt.
Rubrics and verifier	Check whether the scoring criteria can evaluate the final deliverable.	Update rubrics after prompt changes so they match the final task.

Important

The prompt on the left is the one that matters. The prompt on the right is usually a seed or duplicate and can normally be ignored. You can also ignore fields such as “CDS starting file URL to use in the initializer config JSON” and “Verifier template” unless you are troubleshooting a technical issue.

2.2 Source Input files

You will have to source input files if you receive a blank task. You may also add or replace any pre-seeded input files in prefilled tasks.

Download the **initial files** and the **initial output files** (if any) from the Initial Files Upload section and review them. If you are unable to download them by clicking the provided links, copy and paste the links directly into your browser. Save the files locally before continuing, if applicable.

Initial Files Upload

This section is for you to upload the initial supporting files used to create a prompt

Initial Files you are Provided

Starting File names(s): ["dis-nfix-wbd20250628_merged.pdf"]

Locally downloadable starting file URL(s): ["<https://scale-cds-public-us-west-2.s3.amazonaws.com/68a7f8dc133a30ee7f7494de/vIRTa6V6jIA1tvQ>"]

Guidance

1. If you are provided with starting files, move past this step. This means this task already has a prompt that you can improve upon with the provided starting Files. However, if you need to swap out/add to the starting files to make the prompt satisfactory for this project, then please come back here to add the new material
2. If no initial Files were provided, please add cc compliant, high quality starting files that you will use to craft a good prompt. See the instructions for further guidelines on good starting files.

DO NOT UPLOAD GROUND TRUTH FILES HERE



Attachments 0

You may see pre-seeded input files in prefilled tasks. Copy the URL(s) and paste into your browser to view them. You may edit or replace the pre-seeded initial files.

If you would like to upload new input files, click the 'upload attachment' button.

Initial File Check

Input file requirements:

- Minimum of 3 input files required
 - **✗** Input files cannot all be CSV or XLSX file inputs.
- Publicly accessible (no paywalls).
- Open source CC compliant (free for use, fully public license - NO ATTRIBUTION).
- In English.
- Less than 30MB total.
- Relevant to the chosen domain of your task.
- Avoid PII.
- **Stable source:** All tasks must be based on historical or static data. Verify that the data source will not change.

❌ Examples of unstable data sources— REPLACE the files.

- Live stock prices or financial tickers.
- Real-time sports scores.
- Current weather data.
- Live inventory or pricing pages.
- Any page that auto-refreshes or shows 'as of today' figures.

✅ Examples of stable data sources—ACCEPTABLE.

- Historical quarterly revenue reports
- Wikipedia tables for completed events
- Government datasets with fixed publication dates
- Company filings (10-K, 10-Q) with specific period dates
- Archived reports/papers
- **Note: Try to opt out of or modify the arXiv prompt for historic/archived papers.** When working with the arXiv ID, it is possible to retrieve the paper version through the submission history; if this field does not appear, it means it is v1. There are a few ways to use it: the first is <https://arxiv.org/abs/{arxivID}v1> to fetch the version, or you can enter {arxivID}v1 directly into the search engine to access the versioned page.

Acceptable input file formats:

If you have questions on your file format reach out to the QMs

.pdf, .xlsx, .csv, .pptx, .docx, .zip, .jpeg, .png, .txt, .blend are some of the examples for most of the tasks. If an app requires another use that can also be acceptable

2.3 Rewrite the Prompt

Create/Revise the prompt so it reads like a realistic workplace request while remaining specific, feasible, and verifiable. The goal is not to make the prompt longer. The goal is to make it clearer, more natural, and harder in meaningful ways.

A. Required prompt details

Every final prompt must include the following information in natural language:

- ☐ The exact GTF filename or filenames, including the correct extension.
- ☐ The instruction that all output files must be saved on the Desktop.
- ☐ The expected deliverable type, docx .pptx .pdf .zip .blend .tsn .escn .tres .godot .gd .scn .res, etc.
- ☐ The core work the agent needs to complete: analysis, editing, formatting, transformation, research, correction, or creation.
- ☐ Any essential constraints needed to make the task unambiguous and verifiable.

Example of a natural requirement

"I'm working on a presentation for work. Can you review the attached quarterly sales workbook and prepare a leadership-ready summary. Add a chart showing regional revenue by quarter, clean up the print layout so it fits on one page, and save the completed file as `regional\_sales\_summary.xlsx` on the Desktop."

B. Prompt quality standards

Standard	What it means
Clear	A reasonable agent or attempter should not be able to interpret the task in multiple valid ways.
Natural	The prompt should sound like a real request from a manager, client, teammate, or stakeholder.
Complex enough	The task should require meaningful professional judgment, multi-step execution, careful formatting, or cross-file reasoning.
Feasible	The task must be possible in the Ubuntu VM using the provided files and supported tools.
Verifiable	The output must be something the rubric and verifier can compare against the ground truth.
Aligned	The prompt, GTF files, rubrics, and verifier must describe the same final task.

C. What to avoid

Avoid	Why it hurts quality	Better approach
Rubric-style language	It makes the prompt sound like evaluation instructions instead of a real request.	Write as a natural delegation of work.
Over-prescribed algorithms	It turns the task into mechanical rule-following and often reveals the verifier logic.	Give the business goal and necessary constraints, not every internal step.
Artificial personas	"You are a finance manager" often feels decorative and fake.	Use realistic context only when it changes the work product.
Excessive "must / only / exactly" wording	It makes the prompt adversarial and machine-generated.	Reserve strict wording for filenames, save locations, and essential constraints.
Exact section-title checklists	It reduces realism and can overfit the output to a template.	Ask for the content that needs to be covered, and specify structure only when professionally necessary.
Robotic file references	Phrases like "extract data from file A and file B" can feel unnatural.	Refer to the attached workbook, report, deck, image, or dataset in context.
Putting the answer in the prompt	If the prompt states the values, the conclusion, or the exact sentences the deliverable is scored on, the agent can copy them straight into its output and earn credit without doing the work.	Give the task, context, and constraints. Let the agent derive the graded content from the input files.

D. Add meaningful GUI interaction

GUI interaction is a critical part of this project. Many prompts can be solved entirely through code or command-line tools, especially spreadsheet, document, and presentation tasks. A strong Orangutan prompt should include work that benefits from actually opening the file, viewing the rendered output, navigating a browser, adjusting visual layout, or using an application interface.

Core principle

Do not add GUI interaction just for the sake of clicking. Add it when the UI changes what the agent can observe or do: rendered layout, visual overlap, page breaks, chart positioning, animations, comments, tracked changes, browser navigation, image alignment, or app-specific formatting behavior.

GUI pattern	Why it works	Prompt language you can adapt
Rendered layout checks	The agent has to visually inspect the opened file because scripts often cannot reliably determine whether text overlaps, fits, or looks professional after rendering.	

GUI pattern	Why it works	Prompt language you can adapt
Print/page setup	Spreadsheet and document print settings are often easier and more reliable to verify through the GUI than through raw file edits.	"Use page-break or print-preview view to make sure the final sheet prints cleanly on one page."
Visual reference matching	The target is defined by appearance, spacing, proportions, or a screenshot, which pushes the agent to compare visually.	"Recreate the attached reference layout as closely as possible, including spacing, relative element sizes, and placement."
Interactive review features	Comments, tracked changes, animation previews, and slideshow playback are inherently UI-oriented.	"Review the inline comments and tracked changes, then accept the approved edits and resolve the comments."
Browser navigation	Chrome tasks require the agent to move through a website, identify a stable historical/static value, and bring it into the deliverable.	"Navigate to the specified public page and add the listed historical note to the report."
Iterative visual calibration	Some visual decisions require observation and adjustment rather than a single formula.	"Adjust the opacity/axis bounds/label position until the result is readable and visually balanced."

Bad GUI addition

"Open the file and click around before saving."

This does not add meaningful difficulty and is not tied to the deliverable.

Good GUI addition

"Open the presentation and adjust the translated German text so it fits cleanly in the existing text box without overlapping the chart. Preview the slide before saving."

This is meaningful because the agent has to inspect the rendered slide and correct a visual layout problem.

PowerPoint-specific note

Strong presentation tasks often involve visual reasoning, not just content editing. Good examples include checking that a footnote fits inside a defined footer zone, fixing overlapping semi-transparent shapes, matching a screenshot or mockup, adjusting a gradient by visual comparison, or previewing animation timing in slideshow mode.

Use this GUI strategy when...	Avoid it when...
The task's success depends on visual rendering, layout, interactivity, browser navigation, comments, animations, or app-specific behavior.	The GUI step is unrelated to the deliverable or could be removed without changing the task outcome.
The GUI requirement makes the task more realistic for the profession, such as reviewing a slide deck, preparing a print-ready workbook, or editing a marked-up document.	The prompt merely says to open a file, click, or view something without requiring an observable correction or decision.
The expected result can still be reflected in the GTF file and rubric.	The GUI action cannot be verified in the output artifact or through the task's evaluator.

E. Tasks using Blender App

If the prompt requires work done in the Blender app, please design the task with a **.blend output** file instead of asking for screenshots of the 3D structure in a pdf, docx, etc.

You may want to include a .blend input file and ask the agent to make modifications to the input file and provide the resulting .blend file as an output.

Prompt rewrite examples: making GUI interaction meaningful

Weak / artificial	Stronger / natural
"Use the attached PowerPoint to create exactly 5 slides and save it."	"Use the attached PowerPoint to create a 5-slide client update. On the timeline slide, adjust the callouts so none overlap when viewed in slideshow mode, then save the final deck as `client_update.pptx` on the Desktop."
Weak / artificial	Stronger / natural
"Create a chart from the data."	"Create a chart from the data, place the legend on the left, and resize the plot area so the longest label remains readable in the spreadsheet view."
Weak / artificial	Stronger / natural
"Summarize this document."	"Review the marked-up document, resolve the comments that request wording changes, accept the approved tracked edits, and save a clean final copy as `policy_summary.docx` on the Desktop."

2.4 Validate and Finalize the Prompt

After creating/editing the prompt, confirm that the task is safe, feasible, and internally consistent. A clean prompt is not enough if the GTF files, rubrics, or verifier still reflect the old version of the task.

A. PII and safety check

Personally Identifiable Information must not appear anywhere in the task. Check the prompt, starting files, GTF files, rubrics, and other task content.

PII type	Examples
Real non-famous people	Names of private individuals that can be verified online.
Contact information	Full addresses, phone numbers, email addresses, or similar personal identifiers.
Financial information	Banking details, account numbers, payment information, or similar sensitive data.
Other sensitive personal details	Any private personal data that is unnecessary for the task.

PII rule

If PII appears in the prompt, rubric, or GTF files, remove or generalize it.

B. Final prompt checklist

Before moving on, read the final prompt aloud and confirm the following:

- ☐ The prompt sounds like a realistic workplace request.
- ☐ The prompt is clear enough so anyone reading the prompt and the initial files can replicate the GTF without having to guess things.

- ☐ The exact output filename or filenames are stated clearly.
- ☐ The prompt says that all output files must be saved on the Desktop.
- ☐ The task can be completed in the VM using only installed software (LibreOffice, Chrome, standard tools) and the provided files — no unavailable applications.
- ☐ There is only one reasonable interpretation of the request.
- ☐ The prompt does not contain PII.
- ☐ The GUI interaction, if included, is meaningful and connected to the deliverable.
- ☐ The GTF files match the final version of the prompt.
- ☐ The rubrics evaluate the final version of the prompt.
- ☐ The verifier checks the correct GTF files.

Quick Reference: Prompt Editing Decision Flow

Question	If yes	If no
Does the task match the assigned domain?	Continue.	Skip the task if it is not in your domain.
Are the source files usable, stable, supported, and PII-free?	Continue.	Replace the input files.
Is the prompt clear, natural, and complex enough?	Keep or lightly polish it.	Rewrite the prompt.
Does the prompt include the exact GTF filenames and Desktop save location?	Continue.	Add those details naturally.
Does the prompt include meaningful GUI interaction where appropriate?	Continue.	Add a realistic GUI requirement if it strengthens the task and remains verifiable.
Does the GTF file satisfy the final prompt?	Continue.	Update the GTF file.
Do the rubrics and verifier match the final prompt and GTF file?	Submit after final checks.	Update them before submitting.

Prompt changed?

If you received a prefilled task and edited the prompt, make sure to provide the information about what you changed and **be descriptive**. If no change is made, click 'No' and continue.

1/1 completed

Prompt Change Needed *

Item	Details / Examples
Prompt Changed?	If you changed ANYTHING in the prompt (and you almost always should), mark as Yes .
Change Description Requirements	Minimum 50 characters . Submissions below this threshold will be flagged for rejection. Must reference specific changes and rationale for each change.
BAD Descriptions	Just minor tweaks to make it sound natural.
GOOD Descriptions	"Rewrote seed prompt for naturalness. Changed target from all slides to slides 2-5. Removed infeasible Core Request (insert watermark) since LibreOffice Impress does not support this natively. Added context backstory about the teacher preparing class materials. "Preserved all 4 Core Requests. Adjusted slide numbers to match actual file (skeleton said slide 10, file only has 8 slides). Added backstory about quarterly report preparation."

☒ Yes
☐ No

Why was the prompt Changed
Please describe the changes you made to the prompt and why.

Discard changes

Save

3. Skip Asset Upload & Rubric for now

When you reach this step, you do not need to do anything yet.

Continue to the VM step first, complete the required work there, and then return to the Asset & Rubric sections.


- If you cannot skip due to an empty rubric, feel free to add some placeholders in the rubric section to move past this step for now. Please note that you have to come back and rebuild the rubric later.
-

4. Computer Use Step. VM.

About the Virtual Machine (VM)


For now, we are working with **Ubuntu**, which means you'll be using a **Linux** environment.

For those who are not programmers, a **virtual machine (VM)** is essentially a computer running inside another computer. It gives you a separate environment where you can safely open, edit, and save files without affecting your own system.



4.1 Task Initializer

The purpose of the Initializer is to upload the starter files to the VM (Virtual Machine). To do that, the JSON (the text in the input box) needs to be formatted correctly.



For each starter file, there needs to be a **URL** (the link to the file) and a **path** (the destination in the VM).

The initializer may already be embedded with some initial files if you have a prefilled task, however, you must replace the URL and filenames with the correct set of input files associated to your prompt.

If you intend to use the pre-seeded input files: Copy and paste the filenames and URLs from the <i>'Starting File URLs to be used in Task Initializer'</i> section into the initializer JSON code.	Initial Files Upload This section is for you to upload the initial supporting files used to create a prompt Initial Files you are Provided Starting File names(s): ["Rosette 1 Measuremnts.xlsx","Rosette 3 Measuremnts.xlsx","Fatigue_Assessment_Maintenance_Guide.pdf"] Locally downloadable starting file URL(s): ["https://scale-cds-public-us-west-2.s3.amazonaws.com/68a7f8dc133a30ee7f7494de/2jzDCe2M94MOES","https://scale-cds-public-us-west-2.s3.amazonaws.com/68a7f8dc133a30ee7f7494de/YakIT22G_VhmPSa","https://scale-cds-public-us-west-2.s3.amazonaws.com/68a7f8dc133a30ee7f7494de/Ex6FUQdy-XX8_hf"] Starting File URLs to be used in Task Initializer: ["scale-cds://68a7f8dc133a30ee7f7494de/2jzDCe2M94MOES#s3/scale-cds-public-us-west-2","scale-cds://68a7f8dc133a30ee7f7494de/YakIT22G_VhmPSa#s3/scale-cds-public-us-west-2","scale-cds://68a7f8dc133a30ee7f7494de/Ex6FUQdy-XX8_hf#s3/scale-cds-public-us-west-2"]
If you edited or replaced any pre-seeded input files, OR if you uploaded additional new files:	Upload the edited/replacement/new file in the initial files upload section at the beginning of the task (not to be confused with the Asset Upload section). See the instructions for further guidelines on g DO NOT UPLOAD GROUND TRUTH FILES HERE Upload attachment Attachments 0 Copy the link to your uploaded file by clicking on the link symbol to the right of the file name: Replace the old link in the JSON with this new link, and make sure the file name is correct - if you changed the file name, you need to change it in the "Path" section: "path": "/home/docker/Desktop/your_updated_file_name.txt
Initializer JSON explained	
1. Look for "url": then place the input file URL inside "replace with input file URL"	<pre>{ "config": [{ "type": "download", "parameters": { "files": [{ "url": "replace with input file URL", "path": "/home/docker/Desktop/input filename placeholder.docx" }] } }] }</pre>
2. Next, replace the filename associated with the input file URL at the end of "path": "/home/docker/Desktop/input filename placeholder.docx"	<pre> "url": "replace with input file URL", "path": "/home/docker/Desktop/input filename placeholder.docx" }] } }] }</pre>
3. If you see any files with "_output" in the filename, remove that JSON block.	<pre> "url": "replace with input file URL", "path": "/home/docker/Desktop/output filename placeholder.docx" }] } }] }</pre>
4. If the number of files is below OR exceeds the provided entries, delete the	<pre> "url": "replace with input file URL", "path": "/home/docker/Desktop/output filename placeholder.docx" }] } }] }</pre>

extra JSON block OR copy and paste the following JSON block on a new line after the final set of curly braces.

1 JSON block = 1 highlight block in yellow, orange, or pink.

The green parts should remain exactly the same, and your inserted part should look like one of the middle sections - if you insert it at the very end, please note that there should not be any ending comma after the last }, but if you insert the block in the beginning or in the middle, it should end with a comma.

```
    "path": "/home/docker/Desktop/input filename
placeholder.docx"
  },
  {
    "url": "replace with input file URL",
    "path": "/home/docker/Desktop/input filename
placeholder.docx"
  }
]
},
{
  "evaluator": {}
}
```

Example Initializer JSON code in full:

```
{
  "config": [
    {
      "type": "download",
      "parameters": {
        "files": [
          {
            "url":
"scale-cds://68a7f8dc133a30ee7f7494de/D01f5hVPHj10esq#s3/scale-cds-public-us-west-2",
            "path": "/home/docker/Desktop/PAPER_FIC.pdf"
          },
          {
            "url": "replace with input file URL",
            "path": "/home/docker/Desktop/input filename placeholder.docx"
          },
          {
            "url": "replace with input file URL",
            "path": "/home/docker/Desktop/input filename placeholder.docx"
          },
          {
            "url": "replace with input file URL",
            "path": "/home/docker/Desktop/input filename placeholder.docx"
          },
          {
            "url": "replace with input file URL",
            "path": "/home/docker/Desktop/input filename placeholder.docx"
          }
        ]
      }
    }
  ]
}
```




```
}  
  }  
}  
],  
"evaluator": {}  
}
```

5. When all input file URLs and filenames are placed into the initializer, click **"Execute Task Initializer."** This will load the 'input files' of the prompt into the VM!

If contributors would like to use the initial output files (pre-seeded output files) but they are not in the VM, please transfer them to VM yourself (e.g., via filebin, Google Drive, etc).

Execute Task Initializer



⚠️⚠️ **Your GTF must be reproducible from the input files.** Every graded value in your GTF has to be something an agent could derive using only the inputs it receives. Before you finish, trace each key value back to a specific input file. If you can't point to where a value comes from, the agent can't either — and a correct

agent that can't match your GTF means the task fails.

- **Joins and merges:** state the number of records the operation should produce, and verify your GTF has exactly that. A join that silently adds or drops rows is the most common way a GTF stops matching the inputs.
- **Looked-up constants:** if a value comes from a spec, standard, or reference, source it — don't hardcode a number from memory.
- **If you edited or replaced the input files:** rebuild the GTF from the *new* inputs. Building it from the original seeded output means the GTF describes data the agent never receives.

Everything is connected: the prompt, the rubric, and the GTF files should all be consistent with one another and represent the same final ground truth solution.

4.2 Create GTF Files



Workflow: Work Inside the VM



Creating the GTF Files

You have **two options** when creating the ground truth file:

Option 1: Create It From Scratch

- Use the **initial file(s)** as the source of truth.
- Use the **initial output file (if any)** as a reference.
- Build a completely new ground truth file (GTF file) that correctly satisfies the prompt requirements in VM.

Option 2: Review and Edit the Existing Initial Output File (if any)

- Compare the **initial output file** against the **input file(s)** and the prompt.
- Make any necessary edits to ensure the file is completely accurate and represents the correct solution in the VM.



When you finish, download the completed file to your local machine

- You may use Filebin or similar sites to download files to your local machine.

Which files can the agent create?

Your GTF files need to be in a format that the agent can reproduce. At the moment, we know that the model can create files with these extensions:

`.docx .pptx .pdf .zip .blend .tscn .escn .tres .godot .gd .scn .res`

Extensions which are not in this list, such as .txt files or videos, are evaluated on a best-effort basis by the agent judge.

✗ Not Recommended

- Create/Edit the GTF files entirely on your local machine.

⚠ **Note:** Working directly in the VM helps avoid compatibility issues and file transfer problems

⚠ Your GTF must derive from the provided input files. Every graded value in the GTF has to be reproducible by an agent working only from the inputs and the prompt it receives. If your GTF pulls from a different dataset, uses different IDs, or contains figures the agent never gets, a correct agent cannot match it — and the task fails regardless of how good the GTF looks on its own.

Before finishing, trace each key value in your GTF back to a specific input file. If you can't point to where it comes from, the agent can't either.

If the VM does not contain your GTF files on the desktop, the verifier will **NOT** return a 1.

4.3 Upload GTF Files to the Asset Section

Once you have created the GTF files, download them from the VM to your own local machine and return to the previous **asset upload** step.



Uploading the Final GTF Files to the Asset Section

Click "Upload attachment" and when uploading the files:

- Make sure the filename exactly matches the GTF filename specified earlier in the prompt.

FORMAT RULES (LibreOffice on Ubuntu):

- LibreOffice Impress tasks: save as ``.pptx`` (NOT ``.odp``)
- LibreOffice Writer tasks: save as ``.docx`` (NOT ``.odt``)
- LibreOffice Calc tasks: save as ``.xlsx`` (NOT ``.ods``)

NOTE: DO NOT CONVERT PDFs to DOCX. Just select 'all files' only

 Upload attachment

Attachments 0

5. Rubric Section

New! [Rubric Examples doc](#) is now available!

5.1 Overview

After you have the GTF files created in the VM and uploaded to the Asset Section, come back to create the Rubrics. A rubric is a **checklist of criteria** that defines an ideal response to a user prompt.

Well-written rubrics are important because **they are used to evaluate model responses**. A judge, that is, another human or an LLM, will check your rubric criteria against a model response and assess, criterion by criterion, if the response matches those criteria.

You can think of a rubric as follows:

- A **complete, non-redundant**, and **accurate** set of **atomic** truth conditions specifying under which conditions a response to a given prompt is a perfect response. Basically, a **blueprint** for a perfect response to the prompt you wrote.
- Each criterion is an item on the rubric that defines a single, discrete condition needed for the ideal response.

- The LLM grader will NOT have access to the prompt and input files. The rubric must be written in a way that only references the GTF files and the agent response.
- You must produce weighted criteria that a judge model can use to score the agent's output against the GT.

The important shift: the GTF file is the scoring anchor. **The rubric should tell the judge what dimensions to compare, not restate the answer key or reward superficial presence.**

Important note: Start from the prompt, not the gold file. Before writing or editing any criteria, list every discrete requirement in the prompt — each deliverable, each piece of content, each constraint, each formatting or structural ask. Write one criterion per requirement. Then mark each requirement as either:

- **Anchored** — there is one correct answer, so the criterion compares against the expected file; or
- **Open** — the prompt allows multiple valid interpretations (implicit requests, open-ended recommendations, "choose an appropriate X"), so the criterion does *not* anchor to the expected file.

Building the rubric by describing the gold file you produced is the most common source of failure: it leaves prompt requirements uncovered (nothing checks them) and over-anchors open-ended elements to your one specific choice (penalizing other valid outputs). Walking the prompt first prevents both.

NEW: Copilot

The rubric section now contains copilot checks that work better than linters that provide guidelines on what is missing or wrong in your rubric and fulfilling that should result in a good rubric. The main points covered by copilot are, coverage, categorization, overfitting, framing, accuracy and atomicity. Be sure that you read the feedback provided by copilot to ensure your task has a good chance of getting a good score in the review layer

Rubric Generator

 [Rubric Generator Walk-through video](#)

You will have to paste in the (1)prompt, (2)Input Files, and the (3)GTF files from the Asset Upload section.



Password: scaleai2026

1. Paste the prompt
2. Paste the Input file URLs and filenames from the initial files section
3. Paste the GTF file URLs and filenames from the Asset Upload section
4. Enter password:
scaleai2026
5. Click 'Generate Rubric'
6. Download or copy the auto-generated rubric to a separate document to save it before clicking 'Next'.
7. Paste the generated rubric back into the task.

a. Please inspect the auto-generated rubric and make any necessary edits.

b. 'total weight = 100' is NOT a requirement. You may add/delete criteria and/or adjust the weights as long as the weight share (%) matches the guidelines stated in section [5.2 Scoring Model](#).

⚠ You will manually copy & paste each criterion into the rubric builder. Clicking 'save' at the rubric



generator step will NOT autofill the criteria into the builder.



Rubrics are auto-generated based on the GTF files.

Click on the generate button to generate the rubrics

Accuracy & Completeness Check: You must review, edit, and improve the rubric accordingly.

- It is a good practice to always specify the agent output file name and reference the GTF file as the 'expected file' in the rubric
 - E.g., "The 'Experiment_Results.docx' states the same value of ibuprofen mentioned in the corresponding expected file."

No need to mention file location ('save to desktop') in the rubric.

⚠ Remember: If you modify the prompt, you must also update the rubric accordingly. The rubric should always align with the final version of the prompt and evaluate the requirements that are actually being asked.

Rubric coverage and Length

There should be **10-30 rubric items** in total

Rubric Must Always Be:

Atomic:

Bad:

The idea is for each criterion to evaluate only one element in the output files and one categorization of this. An element should be considered a major segment or item in an output file. This can be the contents of a section, a table, a chart, an image, etc. For example, if there is a table that includes multiple items, there only needs to be one criteria that evaluates the correctness of the whole table. There can also be another criteria that evaluates the visual component of the table. Each rubric criterion should test one element in one category(visual,correctness, format gate).	- "Table 1 in US_report.pdf and Table 2 in UK_reports.pdf are semantically equivalent to the corresponding tables in the corresponding expected files" - "The data in Table 1 in US_Reports.pdf is semantically equivalent to the data in Table 1 of the corresponding expected File and is also similar in layout." Good: - "Table 1 in US_report.pdf is semantically equivalent to the corresponding table in the corresponding expected file." - "Table 1 in US_report.pdf contains the equivalent GDP values, HDI values, and GDP per capita values to the corresponding table in the corresponding expected file."
Self-Contained: The evaluator must not be required to retrieve extra content to assess the criterion besides the content of the GTF or the prompt. So you must phrase your criteria in such a way that the item can be graded without any external info. The information can also be included in the input file, whether it is order, a specific number or some layout	✗ "States the correct dosage of ibuprofen." ✗ "States the dosage to be twice the dosage mentioned in `starting file A`" ✓ "The `Experiment_Results.docx` states the same value of ibuprofen mentioned in the corresponding expected file" • Note: `Experiment_Results.docx` = the file that the agent creates; expected file = the Ground Truth files (GTF files) that the attempter provides in the asset section
Written with Positive Wording: Unless the prompt explicitly asks to omit something, the rubric items should always be written in terms of an element being present.	✗ "Fails to create a column named "Sale Prices." ✓ "The 'Experimental Results' table in the `Experiment_Results.docx` creates a column named 'Sale Prices'."

Always start with the simple present tense (e.g., Shows the..., includes a..., provides a..., displays X as..., etc.)	✓ “Experiment_Results.docx` states the same value of ibuprofen mentioned in the corresponding expected file.” ✗ “The response must state the same value of ibuprofen mentioned in the expected file.”
Objective, Verifiable & Factually Correct: Target values and claims that can be fact-checked.	
Comprehensive: Includes 100% of the essential criteria.	
Relevant: No unnecessary items.	
Formatting: When citing variables, file names, formulas, or column/row names, enclose the items in simple backticks for better formatting.	
Specificity: Criteria should state precisely what is expected. They should be binary (true or false). Avoid vague descriptions (e.g., "the response must be accurate"): provide more detail to make the criteria as objectively gradable as possible.	
GT-anchored: every accuracy-based criterion must be evaluable by comparing the agent output against the reference artifact or input files. ⚠ Visual Criteria are on a case-by-case basis: if there are multiple correct interpretations of the prompt request, it shouldn't be anchored to the GTF (expected file); If there is only 1 correct presentation/interpretation of the request, it should be anchored to the GTF (expected file). Implicit requests, or situations where there can be multiple correct interpretations of visual elements, do not need to anchor to the GTF (expected file) because doing so will penalize the agent output for a different but equally correct output. Avoid having an overly restrictive rubric. ⚠ Cross file consistency checks do not need to be GT-anchored. Note that they don't always belong to the correctness category. • Example 1: Prompt: “The Report.pdf should have the same performance metrics table as the metrics.pdf” ○ Rubric “Report.pdf contains the same performance metrics table with the same metric columns and values as the table shown in metrics.pdf” - Correctness • Example 2: Prompt: “Pick a light blue color for the background of Metrics.pptx and use the same color in Report.pptx” ○ Rubric “Report.pptx uses the same background color as the background color in Metrics.pptx” - Visual	

● Example 3: Prompt “Write up a summary of the [XYZ] findings in file.docx, and then create a print-ready pdf of the docx, named poster.pdf” ○ Rubric “poster.pdf presents semantically equivalent findings as the findings reported in file.docx” - correctness (<i>cross file consistency checking if the pdf has the same findings as the docx</i>); ● Example 4: Prompt “Group the complaints into appropriate themes in complaints.csv. Make one summary slide per theme in themes.pptx, include complaint examples for each.” ○ Rubric “Each complaint example shown on a theme slide in themes.pptx is assigned to that same theme in complaints.csv.” - correctness
Correctness-focused: the majority of the score should test the substance of the deliverable, not whether a file or section exists.
Tolerance-aware: allow appropriate numeric rounding, semantic equivalence in prose, and visually comparable layouts.
Exhaustive where outputs are enumerable: cover all countable entries in the GT, not a small spot-check.
Consistency-aware: verify that repeated KPIs, charts, tables, and narrative claims agree across files and pages.
Non-duplicative: score each factual claim or output dimension once.

5.2 Scoring Model

The format confirms that the agent attempted the task. Correctness determines whether the agent succeeded. A plausible-looking deliverable should not earn a high score if the numbers, transformations, classifications, or conclusions are wrong.

NOTE: For Design and Multimedia tasks, the weight share is different, there is not a clear guideline of exact numbers but you can find an estimate in the next table

Scoring area	Target weight share	How to use it	Example criteria
Format Gates (MUST-PASS)/ (Regular)	$\leq 20\%$	For Must pass: Use for file existence and file type only. <i>Counts and non-visual structural checks (page count, slide count, folder structure, document length e.g. word count) are Format Gate, but not Must Pass — they measure whether the work is right, not whether a file was produced.</i>	"Output includes a file named `X.xlsx` with the .xlsx extension."
Correctness (Regular)	$\geq 50\%$ Note: 20-30% for Design & Creative domain	Use for computed values, data transformations, reasoning, classifications, recommendations, and chart data.	"All computed values in `results.pdf` match the values in the corresponding expected file within $\pm x\%$ "; "All 44 waste-route classifications in `example.csv` match the values in the corresponding expected file"; "The recommendation in section 3 reaches the same semantic conclusion as the expected file."

Visual (Regular)	20-30% Note: >=50% for Design & Creative domain	Use for charts, rounding, typography, style document layout, slide presentation, formatted documents, GUI outputs, section and element ordering(when reasoning is not involved) Implicit requirements do not have to anchor to the GTF (expected file). Explicit requirements, where there is only 1 correct interpretation/answer, should anchor to the GTF (expected file).	"The chart axes, scales, and series in `chart_A.jpg` are well spaced and legible." ⚠ If the prompt explicitly requests a specific format/color/layout: • `Hello.pptx` matches the template structure of the expected file/ • `Hello.pptx` matches the structure of the page layout of the expected file. • "The layout polish of `hello.pptx` is professionally comparable to the corresponding expected file." • "The background color and font match the corresponding expected file." Note: Implicit requests, or situations where there can be multiple correct interpretations of visual elements, do not need to anchor to the GTF (expected file) because doing so will penalize the agent output for a different but equally correct output. Avoid having overly restrictive rubric.
------------------------------------	---	--	---

Additional Weight Guidance

Weighting rule: Weights should reflect the importance of a correct deliverable, not how often models fail. Never give a trivial gate the same weight as a core analytical result.

Based on the criteria, make sure that your weights are proportionate to the weights given to other criteria in that scoring area

- **Critical:** These are the core criteria that determine if the agent successfully completed the task. These should generally have the highest weight.
- **Important:** Important but secondary criteria in evaluating the usefulness/accuracy of the agent response.
- **Minor:** These are criteria that should be correct but have only a minor impact on the usefulness of the agent response. These should have the least weight.

‘total weight = 100’ is NOT a requirement.

- You may add/delete criteria and/or adjust the weights as long as the weight share (%) matches the guidelines stated in section [5.2 Scoring Model](#).

5.3 Write Criteria Against the Reference

A good criterion names the dimension to compare and points the judge to the GT artifact. It should not expose a long list of answer values unless those values are explicitly optional.

Instead of writing...	Write this way...	Why
Cell B4 equals \$83.3B.	The revenue figures in the summary table of `Financial_Report.docx` are semantically equivalent to the values in the corresponding expected file.	The judge can compare the expected file; the rubric should avoid answer leakage.
Includes a revenue section.	The analysis section's revenue in `Financial_Report.docx` is semantically equivalent to the section in the corresponding expected file.	Presence of a section does not prove the analysis is correct.
Mentions all waste routes.	For each waste route in `Experimental_Setup.docx`, it includes the corresponding route decision and supporting details in the expected file.	Enumerable outputs should be covered fully, not spot-checked.

Instead of writing...	Write this way...	Why
Has a summary.	The financial summary in `Financial Report.docx` contains values matching those in the corresponding expected file.	Summary existence is not summary correctness.

Note: In the above examples, `file\_name\_placeholder.docx` = the file that the agent creates; corresponding expected file = the GTF files that the attempter provides in the asset section.

5.4 Tolerance Rules

Criteria should be strict about correctness while allowing reasonable variation. State the acceptable comparison mode whenever the prompt permits more than one valid expression.

Output type	Recommended tolerance	Example criterion language
Numeric values	Use a domain-appropriate tolerance such as +/-1%, +/-3%, +/-5%, to account for rounding.	Computed totals and intermediate values match the expected file within the specified tolerance.
Text and prose	Use semantic equivalence unless exact wording is required.	The recommendation conveys the same conclusion and rationale as the expected file, even if phrased differently.
Open-ended methods	Allow defensible alternatives when the prompt does not prescribe one method.	The agent's method produces results aligned with the expected file and is justified by the task constraints.

5.5 What Not To Do

- Do not hard-code literal expected values in criterion titles ("Shows X as 3.89").
 - Instead: Compared to expected via tolerance.
- Do not write existence-only correctness criteria ("Includes a revenue analysis section").
 - Require substantive correctness against the expected.
- Do not give heavy weight to trivial checks (file extension, page count).
- Do not spot-check 7 of 44 enumerable items when you can require all 44.
- Do not penalize valid alternative phrasings, rounding within tolerance, or defensible methodological choices the prompt did not constrain.
- Do not assign all criteria MUST-PASS. MUST-PASS is reserved for format-gate criteria. Everything else is REGULAR.
- Do not assign any criterion a weight above 50 or below 1. Stay strictly in the integer range [1, 50].
- Do not write formatting or visual criteria that a degenerate solution can satisfy. Pair any "fits / no overflow / no overlap" check with a legibility guard.
 - ✗ "Text in the footer does not overflow the box." (passable by shrinking text to unreadable size)
 - ✓ "The footer text in `deck.pptx` fits within the footer zone and remains legible at normal zoom."

- Do not write more than three existence-only criteria per rubric. An existence-only criterion checks that something is merely present (a header, page numbers, a section, a title) without checking that its content is correct. These are weak — they reward superficial presence — so cap them at three and make every other criterion test substance.
 - ✗ "Includes a header," "Includes page numbers," "Includes a classification mark," "Includes headings" (four existence-only criteria)
 - ✓ Keep at most three such checks; convert the rest into content-correctness criteria (e.g. "The header states the correct document title and date matching the expected file")
- Do not write value criteria that only check whether a value *appears* in the file. Bind each value to its correct row, column, or label, so a response can't pass by dumping every candidate value.
 - ✗ "The net revenue figure appears in `summary.xlsx`."
 - ✓ "The net revenue figure in the Q3 row of the summary table in `summary.xlsx` is semantically equivalent to the corresponding expected file."

(This is distinct from atomicity: a criterion can test one value (atomic) yet still fail to bind it to a location, letting a value-dump pass.)

5.5 Copilot checks

Recently a new included copilot checks in task to help with quality and make the tasks easier to be completed.

What is copilot?

Copilot is a set of LLM-based, in-task quality checks, the revamped version of linters. Each check is configured with:

- An evaluator agent + a requirement/definition (the rules the agent judges against)
- A set of selected steps and data fields it reads from the task.

When you work the step the check is attached to, the agent runs and returns guidance in the Copilot panel rather than a simple pass/fail string.

Where has copilot been implemented?

NOTE: All of the descriptions and specifics have already been detailed in the previous sections, this is just an explanation of the copilot checks

1. **Prompt quality:** This will check if a prompt is realistic, complex, feasible, verifiable and has meaningful GUI included in it.
2. **GTF:** This will confirm if the generated GTF fulfills the prompt request without missing anything. It enumerates every discrete ask in the prompt and for each, it

will provide a result of MET, MISSING, or UNVERIFIABLE depending on the GTF content.

3. **Weight protocol:** It will confirm the individual weight of each criteria and the weight share of each category
4. **Overfitting:** This will confirm if any criterion is overfit. Remember that a criterion is overfit only when its wording imposes a concrete, unjustified constraint that would reject a materially different but prompt-compliant solution.
5. **Coverage:** This evaluates that all the requirements from the prompt are included in the rubric, this helps determine whether the rubric would prevent a clearly wrong response from passing on every explicit requirement that matters to the task's intended outcome
6. **Categorization:** Analyzes if every criterion used in the rubric is correctly categorized. Ignore weights.
7. **Rubric:** Evaluates every other single aspect of the rubric mentioned in this document that has not been already evaluated by the previous checks.
8. **Accuracy:** Checks every criterion for objective factual or logical accuracy.
9. **Verifier:** Checks that the verifier and initializer are configured correctly and have the correct data format

These checks are still a work in progress and we will continue to improve them on a recurrent basis, they have been enabled because they already flag most of the errors, there will be times where they overflag, but the amount of times that that happens is low. Please be sure to not dismiss the flags and read through them before continuing your workflow.

6. Verifier

The purpose of the Verifier is to assess whether your uploaded GTF Files fulfill all the criteria in your rubric. In order to do that, it needs to have:

1. The **destination** (URL) and **path** (file name) to the GTF files that you have uploaded in the Asset Upload section - these are referred to as the **Expected Files**.
2. The **destination** and **path** to the GTF files located in the VM - these are referred to as the **Results**.
3. The rubric

Of these, you need to take care of 1 and 2 - the rubric part of the JSON will not require anything from you; it is taken directly from the Rubric section of the task.

PASTE/REPLACE the 'paths' with the correct URL links & 'dest' with the correct GTF filenames


Which Links/Filename?

The links/filenames corresponding to the files that you:

- Created/modified inside the VM.
- Downloaded to your local machine
- Uploaded in the **Asset Upload** step

Click the URL icon in the Asset Upload section to copy the generated links.

You will need to use those links in the following steps



Verifier JSON explained:

1. This is the beginning of the Verifier JSON - no need to touch the green-highlighted parts.

```
{  
  "config": {},  
  "evaluator": {},  
  "func": "agent_judge_multi",  
}
```

2. 'Expected' section:


The sections highlighted in green should not be edited.

For the part highlighted in yellow, make sure these are the actual names of your GTF files.

For the part in orange, copy the URL from the **asset upload** section for the

```
"expected": [  
  {  
    "dest": "your_gtf_file_1.txt",  
    "path":  
    "scale-cds://65cbc42b32ffab95dd54b864/rh7_7X2H1oJcWFG#s3/scale-cd  
s-public-us-west-2",  
    "type": "cloud_file"  
  }  
]
```

file with the same name and paste it in place of whatever was there before.



If you need to add more files to the 'expected' section, add a new section. 1 Section = the blue or pink highlighted text.

For the last of the Expected sections, do not use a comma after the curly bracket.

```
},  
{  
  "dest": "your_GTF_file_2.txt",  
  "path":  
"scale-cds://65cbc42b32ffab95dd54b864/rH7_7X2H1oJcWFG#s3/scale-cd  
s-public-us-west-2",  
  "type": "cloud_file"  
}  
],
```

3. 'Result' section:

The green parts should not be edited, but if the entire result section is missing from your task, you should copy this entire template (including the green sections at the beginning and end).

Replace the sections in yellow with the names of your GTF files. If you have more files, add more sections.

1 Section = the blue or pink highlighted text.

For the last of the Result sections, do not use a comma after the curly bracket.

```
"result": [  
  {  
    "dest": "your_GTF_file_1.txt",  
    "path": "/home/docker/Desktop/your_GTF_file_1.txt",  
    "type": "vm_file"  
  },  
  {  
    "dest": "your_GTF_file_2.txt",  
    "path": "/home/docker/Desktop/your_GTF_file_2.txt",  
    "type": "vm_file"  
  }  
],
```

⚠ You may encounter a verifier JSON code that is missing the 'expected' and/or 'result' sections. If this is the case, copy and paste the example template into the verifier and replace the placeholders with the correct URL and GTF file names.

You may also encounter a JSON code where the 'result' comes first or the 'expected' section comes first. Please check the JSON carefully.

Example verifier template in full:

```
{
  "config": [],
  "evaluator": {
    "func": "agent_judge_multi",
    "expected": [
      {
        "dest": "Target_Architecture_Playbook.docx",
        "path": "scale-cds://65cbc42b32ffab95dd54b864/rH7_7X2H1oJcWFG#s3/scale-cds-public-us-west-2",
        "type": "cloud_file"
      },
      {
        "dest": "Replace with output filename corresponding to the file URL",
        "path": "replace with output file URL found in the asset section",
        "type": "cloud_file"
      }
    ],
    "result": [
      {
        "dest": "Target_Architecture_Playbook.docx",
        "path": "/home/docker/Desktop/Target_Architecture_Playbook.docx",
        "type": "vm_file"
      },
      {

```

```
"dest": "Replace with output filename",
"path": "/home/docker/Desktop/Replace with output filename",
"type": "vm_file"
}
],
"rubric": {
  "criteria": [
    {
      "annotations": {
        "criteria_type": "MUST-PASS",
      },
      "id": "9176bfc1-64be-4169-997a-e4435c0ac1ed",
      "title": "Includes a CSV file named \"Target_AWS_Resources.csv.\",",
      "weight": 30
    },
    {
      "annotations": {
        "criteria_type": "REGULAR"
      },
      "id": "b23ca120-d452-45d1-b313-c0747757ab3c",
      "title": "Reduces the font size in Target_Architecture_Playbook.docx to make the bottom half of the document blank.",
      "weight": 5
    }
  ]
}
```


Do not make edits in the "Rubric" section of the JSON code

For those who are not familiar with JSON or programming, please make sure to:


- Keep the quotation marks (") at the **beginning** and **end** of each link.
- Make sure each link ends with a comma (,), otherwise it's going to give an error.

Removing the quotation marks or commas may cause the configuration to fail.

5. Click on Execute Verifier

NOTES: If you run the verifier and it returns 1 (which means 100% correct), it does not necessarily mean that everything is correct.

- You still need to check if the rubrics are correct and the GTF file matches the data in the input files and the prompt.
- If the rubrics and the GTF file are equally incorrect, the verifier can still point to a successful verification.



```
Output
1 {
2   "evaluationScore": 1,
3   "message": "agent_judge_multi rubric_only verdict=PASS score=1.00
(rubric_score=1.00, verdict_score=1.00 [unused]): Both files are p
and contain all required data with exact values matching every rub
criterion.\n Rubric breakdown:\n   [must-have]
6d032508-4d32-4d25-9aac-25996fd12187: score=1.00 weight=10.0
contribution=10.00 (DOCX file is present.)\n   [must-have]
fa937323-7c8d-46c1-a3be-a852bcc68e8a: score=1.00 weight=10.0
contribution=10.00 (PDF file is present.)\n   [must-have]
```

If the VM does not contain your GTF files on the desktop, the verifier will **NOT** return a 1.

If the verifier does not return a 1, but you've checked that the file names and URLs are correct in both the VM desktop and verifier, **you may proceed with the task.**

Troubleshoot: Issues with Expiring URLs

You may encounter an issue with the *initializer* and/or *verifier* when doing your task, where the URL converts to a signed URL with an expiration.

Fix:

You can fix this by going back to the Initial File section, where you'll find the field:

"CDS starting file URL to use in the initializer config JSON: [scale...]."

- Copy those URLs and paste them into the initializer. Replace any current URLs that contain the expiring link issue with these URLs.
- Check the Verifier as well and make sure there are no expiring URLs there either.
 - If you find any, copy and paste the correct URLs again from the Asset Upload step as explained before.

Make sure the URL does not contain "expire" in it:

`https://s3.labeling-data.net/scale-cds-public-us-west-2/68a7f8dc133a30ee7f7494de/97pVHDxNIVynrfs?Expires=1780884000&Key-Pair-Id=K3DN68TU3V3L3W&Signature=`

Should look like this, without the **expire** parameter:


`scale-cds://68a7f8dc133a30ee7f7494de/97pVHDxNIVynrfs#s3/scale-cds-public-us-west-2`

⚠ Be aware: If a task contains an expiring URL, it will not work for reviewers. As a result, your task could receive a **lower score** or be sent back for fixes. Always verify that all URLs have been properly updated before submitting.

7. GTF eval check

After saving the verifier, you will see a Quality Confirmation question.

Click "Yes" if everything looks good (URLs and filenames are correct in both initializer and verifier JSON)



A screenshot of a 'Quality Confirmation' dialog box. At the top right, it says '1/1 completed'. The title is 'Quality Confirmation *'. Below the title, it says 'Please double check that you completed this task accurately.' There are two radio buttons: 'Yes' (which is selected) and 'No'.

You will see a GTF eval in the next step. It is designed to help sanity check your GTF files, but it is not a replacement for the normal detailed checks of your own work! Please do not trust the eval 100%.

1. Enter the prompt and place the URLs for both the GTF files & input files.
 - a. You may grab the URLs directly from the initializer and verifier JSON, or from the Asset Uploads and initial files upload section.

Prompt

For the second image, use the second table in "Data.xlsx" to identify the closed heat engine subsystem and highlight it by drawing a black dashed rectangular around the complete subsystem without obscuring any labels or process lines. Keep the remainder of the diagram unchanged and save the result as "Sub system diagram.png".

Before saving, open both edited images to visually confirm that the recolored arrows remain aligned with the original process streams and that the dashed subsystem boundary does not overlap any labels, symbols, or arrows.

Ground Truth Files
The answer-key artifacts to judge (required).

+ Add File

1	scale-cds://65cbc42b32ffab95dd54b864/2_7tr_tAzKMtngZ*s3/scale-cds-public-us-w	Sub system diagram.png	Remove
2	scale-cds://65cbc42b32ffab95dd54b864/bH4EAzJQ-oZgMBR#s3/scale-cds-public-us	Color Coded Diagram.png	Remove

Input / Source Files
Optional source material used to fact-check the ground truth.

+ Add File

1	https://scale-cds-public-us-west-2.s3.amazonaws.com/68a7f8dc133a30ee7f7494de/	Key.png	Remove
2	https://scale-cds-public-us-west-2.s3.amazonaws.com/68a7f8dc133a30ee7f7494de/	Data.xlsx	Remove
3	https://scale-cds-public-us-west-2.s3.amazonaws.com/68a7f8dc133a30ee7f7494de/	Diagram.png	Remove

2. Click "Run Evaluation", you should see the eval result in a couple of minutes.



Please do not trust the eval 100%! You should still carefully check all aspects of your work. The eval is meant to be a sanity check in case there is anything missed.

Evaluation Result

FAIL

7/8 requirements met

The Color Coded Diagram is largely correct: process-flow arrows are recolored per the legend (H₂ streams 1-9 black, N₂+H₂O streams 11-13 green, Helium streams W-X-Y-Z red, Air and mixed-composition streams 10/c treated as Mixed Gas in light blue), and the legend from Key.png is embedded. However, the Sub system diagram has a clear violation: the black dashed rectangle drawn around the closed heat engine subsystem cuts through the 'Burner' text label (top-left of the rectangle) and passes across the '3' stream label near the bottom. The prompt explicitly requires the boundary to be drawn 'without obscuring any labels or process lines' and includes a pre-save visual verification step to catch exactly this. The rectangle's placement therefore fails the stated constraint.

Requirements

Requirement	Status	Evidence
Create two PNG files with the exact names 'Color Coded Diagram.png' and 'Sub system diagram.png'	MET	Both files exist in ground_truth_files/ with the required names
Recolor process-flow arrows in Diagram.png using Table 1 fluids and Key.png legend (H ₂ =black, N ₂ +H ₂ O=green, He=red, Mixed Gas=light blue)	MET	Color Coded Diagram.png: streams 1-9 shown black (Hydrogen), streams 11-13 green (Water,Nitrogen), helium loop W-X-Y-Z red, air and multi-component streams (a/b/c and 10) light blue as Mixed Gas
Treat fluids not exactly matching legend categories as Mixed Gas (Air, Water/Hydrogen/Nitrogen/Oxygen mixture)	MET	Air streams (a, b) and stream 10 (multi-component mixture) rendered light blue in Color Coded Diagram.png
Preserve original layout, labels, and component positions	MET	Color Coded Diagram.png retains all original component positions and labels
Include the legend from Key.png in the Color Coded Diagram	MET	Legend box visible in lower-left of Color Coded Diagram.png with all four categories
Identify closed heat engine subsystem from Table 2 (Compressor W→X, Q Hot HE X→Y, Expander Y→Z, Heat Sink HE Z→W) and draw a black dashed rectangle around it	MET	Sub system diagram.png shows a black dashed rectangle around the right-side helium loop containing the compressor, expander, and two heat exchangers
Dashed rectangle must not obscure or overlap any labels, symbols, or arrows	NOT MET	Sub system diagram.png: the top edge of the dashed rectangle cuts through the 'Burner' text label; the bottom edge passes across the '3' and '2' stream labels near the pump
Keep remainder of the diagram unchanged in the Sub system diagram	MET	Sub system diagram.png shows the original diagram with only the dashed rectangle added

Findings

MAJOR

ground_truth_files/Sub system diagram.png

The black dashed rectangle overlaps the 'Burner' text label (top-left vertical dashed line cuts through the word 'Burner') and overlaps the stream '3' label and stream '2' label area at the bottom, directly violating the prompt's explicit constraint that the rectangle must be drawn 'without obscuring any labels or process lines' and the verification step 'the dashed subsystem boundary does not overlap any labels, symbols, or arrows'

Reference: Prompt paragraph 3 and verification paragraph; ground_truth_files/Sub system diagram.png visual inspection

Agent Run (Removed from Attempts)

Agent Run Requirements

- At the end of the task, please click "Run" in the Agent Run section. It takes approximately **20-30 min** to generate

the agent responses.

The goal is to have an average score $\leq 65\%$ among the 3 agent responses.

If the average score is $> 65\%$, you must revise the rubric/prompt and try again.

- Check if the rubric is accurate and complete
 - Check if there is anything off in the agent output files and add to the rubric
 - Also, make sure to double-check that your GTF is actually correct
- The prompt may be too easy. Revise the prompt to make it more challenging.

2. Click on 'Collected Artifacts' to review the agent output files. If you notice that the agent output files have weird formatting or areas for improvement (accuracy) but the existing rubric is not penalizing these issues (overlaid text, space gap, extra unnecessary dots, etc.):

- You must make adjustments to the rubric based on the agent output files, even if the agent score average is $\leq 65\%$.
- The goal is to create a rubric that serves as a **blueprint for a perfect response** to the prompt you wrote.

If the GTF file has completely different answers compared to the agent output files, double-check the accuracy of your GTF files.

Task Runner

TASK ID

aeh-1782510979597 v3



RUNS (16)

⚠ Your last run is saved and shown below. The task builder has newer changes that aren't part of it – click Run only if you want to start a new run with those changes.

▶ Run

● aeh-17825109795... v3 100% 26m 0s Re-run from step	● aeh-17825109795... v3 77% 24m 0s Re-run from step	● aeh-17825109795... v2 0% 5m 0s Re-run from step	● aeh-17825109795... v2 100% 23m 0s Re-run from step	● aeh-17825109795... v2 5% 20m 0s Re-run from step
--	---	---	--	--

Desktop Trajectory Viewer Verifier Results **Collected Artifacts (3)** Conversations Task Run Context

3 files.

EXPECTED Expected gold (1)

Click to download and view agent output files

RESULT Agent deliverables (1)

⚠ Note:

If you clicked the purple 'Run' button more than 1 time, you will see 3 new agent outputs generated on the far left of the output section, but the previous 3 agent responses will stay. So it may look like there are more than 3 agent outputs.

Example

Description

Agent Produced File

GTF File

The agent-produced file has an awkward and unnecessary space gap after the Thesis Statement.

- Go back to the rubrics section to add a criterion regarding spacing

The agent produced a file with unreadable text overlaid in the tables.

- Go back to the rubric section to add a criterion regarding text formatting

Analytical Crosswalk: U.S. Fiscal Sustainability Indicators

Thesis Statement

Federal net operating cost, the CBO unified deficit, and Z.1 federal liabilities measure distinct phenomena under different accounting bases, measurement rules, and institutional boundaries; therefore they are not interchangeable proxies.

Show or hide document and s

Analytical Crosswalk Appendix

Thesis

These indicators are not interchangeable: accrual cost, budget deficit, public-held debt, and Z.1 liabilities use different bases, timing, measurements, and boundaries, in each answers a different risk question.

Verdict

Verdict: Not interchangeable. This is the same verdict used in reconciliation memo.pdf

7-concept crosswalk matrix

Concept	General analysis	Timing boundary difference	Measurement rule	Institutional boundary	Use and comparability verdict
1. GAO consolidated statements	CBO budget hearings	Accrual-based reporting; costs recorded on a cash basis; excludes Federal Reserve S	Only cash-based or derived metrics; excludes non-cash items	Consolidation entities of the U.S. government; excludes non-federal entities	Financial position and operating cost; not interchangeable with Z.1 debt flows
2. Net operating cost	CBO deficit	Revenues recognized when earned; expenses recorded on modified cash basis	Net cost less tax and other revenues; includes incurred costs, actuarial assumption changes, and other non-cash items	Same consolidated reporting entity as the financial statements	Not a cash borrowing measure; not interchangeable with the CBO deficit or Z.1 debt flows
3. CBO deficit	Net operating cost	Revenues recognized when earned; expenses recorded on modified cash basis	Net cost less tax and other revenues; includes incurred costs, actuarial assumption changes, and other non-cash items	Same consolidated reporting entity as the financial statements	Not a cash borrowing measure; not interchangeable with the CBO deficit or Z.1 debt flows
4. CBO debt held by the public	Z.1 Treasury securities: Federal Reserve	Face value of marketable + nonmarketable Treasuries; excludes Federal Reserve holdings (trust funds)	Face value of marketable + nonmarketable Treasuries; excludes Federal Reserve holdings (trust funds)	Public creditors; excludes intragovernmental holdings	Best measure for budget imbalance and borrowing need, but it omits many accrual cost changes
5. Z.1 federal liabilities	CBO debt and GAO liabilities	Face value of marketable + nonmarketable Treasuries; excludes Federal Reserve holdings (trust funds)	Face value of marketable + nonmarketable Treasuries; excludes Federal Reserve holdings (trust funds)	Public creditors; excludes intragovernmental holdings	Best measure for budget imbalance and borrowing need, but it omits many accrual cost changes

Reconciliation Memo: U.S. Fiscal Sustainability Indicators

To: Sovereign Risk Committee | From: Public Sector Financial Economist | Date: 23 June 2026

Executive Verdict

The three indicators are **NOT interchangeable**. The GAO net operating cost is an accrual-based comprehensive measure that recognizes long-dated employee, veteran, and environmental obligations as expense when incurred. The CBO unified deficit is a near-cash measure that records most outlays when paid and revenues when collected, with credit programs on a NPV (FCRA) basis. The Z.1 change in federal liabilities is a statistical, largely market-value flow that follows SNA-style sector boundaries. Using net operating cost as a proxy for the CBO deficit or for Z.1 debt changes would mislead the committee on size, direction, and timing of fiscal stress.

Basis Taxonomy Table

Indicator	Accounting basis	Measurement rule	Institutional boundary
1. GAO consolidated statements	Accrual (FASB SFFAS)	Historical cost + actuarial PV of social insurance and pension obligations; excludes Federal Reserve S	Consolidation entities of the U.S. government; excludes non-federal entities
2. Net operating cost	Accrual (FASAB)	Revenues recognized when earned; expenses recorded on modified cash basis	Same consolidated reporting entity as the financial statements
3. CBO deficit (unified)	Modified cash / FCRA	Outlays when paid; revenues when collected; credit programs on NPV basis	Same consolidated reporting entity as the financial statements
4. CBO debt held by the public	Debt-financing concept	Face value of marketable + nonmarketable Treasuries; excludes Federal Reserve holdings (trust funds)	Public creditors; excludes intragovernmental holdings
5. Z.1 federal liabilities	Statistical (Flow of Funds)	Mostly market value; some at par; net flows by instrument	Federal government sector (excludes monetary authority, GSE)
6. Comparability verdict	Non-equivalent	Each metric answers a different question; difference in measurement rules all differ	Same consolidated reporting entity as the financial statements

Eight Reconciliation Mechanisms (Timing / Measurement / Boundary)

#	Mechanism	Class	Explanation
1	Accrual vs. cash recognition of expenses	Timing	GAO accrues interest, salaries, and grants when incurred; CBO records most outlays when checks clear
2	Actuarial revaluation of social insurance and pension obligations	Measurement	GAO recognizes changes in PV of Social Security, Medicare, military/civilian pensions, and veterans' co
3	Credit-program subsidy estimation (FCRA vs. fair value)	Measurement	CBO records credit programs on FCRA-discounted basis; GAO uses similar estimates with re-estimates
4	Market vs. par value of Treasury debt	Measurement	Z.1 marks marketable Treasuries to market and revalues TIPS for inflation; CBO's debt-held-by-the-publ
5	Trust-fund intragovernmental holdings	Boundary	CBO 'debt held by the public' excludes intragovernmental debt; GAO consolidates it out; Z.1 treats intrag
6	Federal Reserve treatment (SOMA, reverse repo)	Boundary	Z.1 places the monetary authority in a separate sector; GAO excludes the Federal Reserve System from
7	Government-sponsored enterprises and GSEs	Boundary	GAO excludes GSEs and Agency- and GSE-backed MBS pools outside the federal sector; CBO and GAO e
8	Seigniorage, financing accounts, and other means of financing	Boundary	Seigniorage, changes in cash balances, and non-budgetary financing accounts reconcile the unified def

Classification summary: 2 timing (#1, #8), 3 measurement (#2, #3, #4), 3 boundary (#5, #6, #7).

Fiscal Indicator Reconciliation Memo

Executive verdict

Not interchangeable. Federal net operating cost, the CBO budget deficit, CBO debt held by the public, and Z.1 federal liabilities answer related but different questions because their accounting bases, measurement rules, timing conversions, and institutional boundaries differ.

Decision: reject the analyst's assertion that P2024 federal net operating cost can proxy both the CBO budget deficit and the Z.1 change in federal debt. Use the three series as complementary diagnostics rather than substitutes.

Basis taxonomy table

Framework indicator	Accounting basis	Measurement rule	Institutional boundary	Use and comparability verdict
GAO consolidated statements	Federal accrual financial reporting under federal GAAP; most revenues recorded on modified cash basis	Costs, revenues, assets, liabilities, net position, and reconciliations are reported for the government-wide reporting entity	Consolidation entities of the U.S. government; excludes non-federal entities such as the Federal Reserve System and Family MaxProtein Inc. are not consolidated	Best for financial position and operating cost, but the GAO disclaimer limits assurance and prevents direct substitution for budget or Z.1 statistics
Net operating cost	Accrual-based operating result	Net cost less tax and other revenues; includes incurred costs, actuarial assumption changes, and other non-cash items	Same consolidated reporting entity as the financial statements	Not a cash borrowing measure; is not interchangeable with the CBO deficit or Z.1 debt flows
CBO deficit	Federal budget framework based on receipts and outlays	Deficit equals outlays minus receipts; CBO also adjusts some presentations for payment timing shifts	Federal budget accounts and budget enforcement conventions	Best measure for budget imbalance and borrowing need, but it omits many accrual cost changes
CBO debt held by the public	Budget debt concept linked to Treasury borrowing	Treasury debt held by outside investors and other non-federal holders, reported as a stock of fiscal year end	Public creditors; excludes intragovernmental holdings	Best budget-sharing debt burden indicator; not a full balance-sheet liability measure
Z.1 federal liabilities	Financial accounts; flow-of-funds system	Quarterly transactions and levels by sector and instrument; levels are end-of-quarter stocks and may include valuation, source-side, and statistical-discrepancy effects	Federal government sector in the macro-financial accounts	Best macro-financial liability view, but sectoring and instrument valuation differ from budget and GAO reporting
Fair comparability verdict	Cross-framework judgment	No single conversion makes all indicators equivalent	Boundaries do not align across consolidated reporting entities; budget accounts, public holders, and Z.1 sectors	The indicators are related but not interchangeable

Issues Reporting

If you encounter any issues with the Agent Run step (e.g., agent fails to run), please provide feedback in the feedback box at the end of the task.

You may submit the task if the agent run step returns an error. But please indicate this in the agent run issues feedback section.

Collect text inputs from the user

1/2 completed

Agent run Working? *

Did the agent run step work for you?

no

What issue did you encounter *

Please describe the issue you faced on the agent run step. Please note it is normal for it to take up to 30 minutes

Next

! When in doubt, do not submit

If you are unsure about data stability, evaluator coverage, or GTF file accuracy, hold the task and ask in the team channel. Submitting an incorrect task creates a downstream QC load for the entire pipeline.

For questions, post in the Orangutan Encyclopedia chat in the Outlier Community. Do not submit tasks with unresolved evaluator issues.